

8 PQHS 471 Notes Week 8

8.1 Week 8 Day 1

8.1.1 ISLR 8.2

8.1.2 Random Forests

Random forests are a model averaging technique using trees as base models. Bagging is a special case of random forests.

1. Specify B , the number of bootstrap samples, and $m \leq p$, the number of features to consider at each split.
2. For $b = 1, \dots, B$, create a bootstrap sample of the training data, build a maximal-depth tree \hat{r}_b *without pruning*. For every split during tree building, only m randomly selected features are considered.
3. The final RF model is $\hat{r}_{\text{rf}} = \frac{1}{B} \hat{r}_b$ for continuous outcome, or majority vote over \hat{r}_b for categorical outcomes.

When $m = p$ for every split, the method is called **bagging**. In R `randomForest` package, by default, m (`mtry`) is $m = \sqrt{p}$ for classification trees and $m = p/3$ for regression trees, B (`ntree`) is 500. The measure of impurity is MSE for regression trees and Gini index for classification trees.

Notes:

- m is a hyperparameter. Once m is specified, RF is an automatic procedure and no tuning is needed. The choice of m determines the complexity of the final model.
- B is not a tuning parameter because a large value of B will not lead to overfitting.
- This algorithm is non-deterministic.
- RFs are **adaptive nearest-neighbor estimators** with neighborhood-defining features selected adaptively. The value m implicitly determines the definition of neighborhood; the smaller m the larger neighborhood.
- Maximal-depth reduces correlation from one tree to another. Compared to bagging, the additional randomization due to a small $m < p$ results in trees with a higher variance, but are a lot less correlated.
- Maximal-depth can be slow for very large n . In `randomForest()`, one can set `nodesize=` (minimum size of terminal nodes) and `maxnodes=` (maximum number of terminal nodes).
- Majority vote is a type of average.
- Model averaging can be used with any base models besides trees.

Justification for model averaging: Let $\hat{a}_1, \dots, \hat{a}_B$ be estimates of μ that have the same mean μ and variance σ^2 . If their pairwise correlation is $\text{cor}(\hat{a}_i, \hat{a}_j) = \rho$ for any $i \neq j$, then their average $\frac{1}{B} \sum \hat{a}_i$ has mean μ and variance

$$\frac{\sigma^2}{B^2} [B + B(B-1)\rho] = \sigma^2 [\rho + \frac{1}{B}(1-\rho)].$$

We can pick a large B to make $\frac{1}{B}\sigma^2(1-\rho)$ very small. If ρ is very small and σ^2 is not too high, then $\sigma^2\rho$ can be small. Thus, we seek to define a base model (or classifier) that (1) is unbiased, (2) is fast to build (so that B can be large), (3) has very small correlation from one to another and a variance not too large.

`library(randomForest)`

```
Heart = read.csv("Heart.csv", row.names=1) ## first column is row name, not a variable
names(Heart); dim(Heart) ## 303, 14
Heart$Thal = factor(Heart$Thal, c('normal', 'reversible', 'fixed'), ordered=T)
Heart$ChestPain = factor(Heart$ChestPain, levels(Heart$ChestPain), ordered=T)

rf = randomForest(AHD ~ ., data=Heart) ## error due to missing data

apply(is.na(Heart), 2, sum) ## check which variable has missing data
apply(is.na(Heart), 1, sum) ## check which observation has missing data

Heart2 = na.omit(Heart); dim(Heart2) ## 297, 14
set.seed(899); rf1 = randomForest(AHD ~ ., data=Heart2)
```

```
set.seed(899); rf2 = randomForest(AHD ~ ., data=Heart, na.action=na.omit)
all.equal(rf1$predicted, rf2$predicted, check.attributes=F) ## True
```

Now we fit a RF model:

```
rf = randomForest(AHD ~ ., data=Heart2)
rf
names(rf)
rf$mtry; rf$ntree
```

Out-of-bag (OOB) error estimation: RF allows CV-like error estimation without CV. For every tree, those in the bootstrap sample are training data and those left out are test data. Let w_{bi} be the number of copies of observation i is in bootstrap sample b . Let $B_i = \#\{b : w_{bi} = 0\}$ be the number of bootstrap samples not containing observation i . (Note that $E(B_i) = e^{-1}B \approx 0.37B$.) For observation i , let $\hat{r}_{\text{rf}}^{(i)}(x_i) = \frac{1}{B_i} \sum_{b:w_{bi}=0} \hat{r}_b(x_i)$, or majority vote. Then $\text{OOB}_i = L(y_i, \hat{r}_{\text{rf}}^{(i)}(x_i))$, and $\text{OOB} = \frac{1}{n} \sum_i \text{OOB}_i$. When B is sufficiently large, the OOB error estimate is equivalent to leave-one-out cross-validation error.

```
mean(rf$oob.times/rf$ntree); exp(-1)
```

For a model fitted with `randomForest()`, its component `err.rate` contains the overall OOB error and the OOB errors stratified by the outcome categories as the number of trees increases. The function `plot.randomForest()` displays these errors in a plot.

```
plot(rf) ## same as matplot(rf$err.rate, type='l', xlab='trees', ylab='Error')
matplot(rf$err.rate[, 'OOB'], type='l', xlab='trees', ylab='Error') ## overall OOB error only
```

All measures of model performance are based on the OOB samples: `votes` contains the distribution of OOB predictions, `predicted` is the result of majority vote, `confusion` is the corresponding confusion matrix.

```
table(rf$predicted, rf$votes[,2] >= 0.5) ## "predicted" is based on "votes"
rf$confusion ## same as table(Heart2$AHD, rf$predicted)
rf$err.rate[rf$ntree, ]
```

We can make predictions on new data using `predict()`. BUT we should not re-predict the training set because it would give near perfect accuracy. This is called the “apparent error rate” and is known for being wildly optimistic.

```
#rfyhat = predict(rf, testdata)
rfyhat = predict(rf) ## without newdata, this returns rf$predicted
```

```
table(Heart2$AHD, predict(rf))
table(Heart2$AHD, predict(rf, Heart2)) ## This is wrong!
```

Variable importance: To measure the importance of a variable, we record the reduction of impurity (MSE, Gini, etc.) whenever the variable is used to split a node, compute the total reduction for every tree, and then average over all trees.

For a classification model fitted with `randomForest()`, its component `importance` contains `MeanDecreaseGini`. The function `importance()` and `varImpPlot()` display the importance values.

```
rf = randomForest(AHD ~ ., data=Heart2)
importance(rf) ## show rf$importance
varImpPlot(rf) ## same as dotchart(rf$importance[, 'MeanDecreaseGini']) except order
```

The results can vary quite a bit between runs because of the randomness introduced during model fitting: (1) bootstrap sampling, and (2) a random subset of features are considered at each split.

```
varImpPlot(randomForest(AHD ~ ., data=Heart2)) ## repeat a few times
```

Below we evaluate the variation of results across runs.

```
importance.multirun = matrix(,20,13)
for(i in 1:20)
```

```
importance.multirun[i,] = randomForest(AHD ~ ., data=Heart2, ntree=500)$importance
colnames(importance.multirun) = rownames(rf$importance)
idx = order(apply(importance.multirun, 2, median))

par(mar=c(4,6,1,1))
boxplot(importance.multirun[, idx], horizontal=T, las=1, ylim=c(0,20))
#dotchart(importance.multirun[, idx], las=1, ylim=c(0,20))
```

The smaller m the more spread-out of importance over the variables. A small m has some similarity to ridge regression, which tends to share the coefficients evenly among correlated variables.

```
varImpPlot(randomForest(AHD ~ ., data=Heart2, mtry=1))
varImpPlot(randomForest(AHD ~ ., data=Heart2, mtry=3))
varImpPlot(randomForest(AHD ~ ., data=Heart2, mtry=10))
```

Another measure of **variable importance** is the *mean decrease in accuracy*, which is calculated by `randomForest()` when `importance=T` is specified.

```
rf = randomForest(AHD ~ ., data=Heart2, importance=T)
importance(rf) ## different from rf$importance except the last column
varImpPlot(rf)
```

The help page for `importance()` explains what the measure is: For each tree b , the prediction error e_b on the OOB portion of the data is recorded (error rate for classification, MSE for regression). Then for each predictor j , permute it in the OOB portion, calculate the prediction error again as e_{bj} , and $d_{bj} = e_{bj} - e_b$, which is the amount of additional error caused by permuting predictor j (effectively making it noninformative). The mean of d_{bj} over all trees is stored in the component `importance` while the standard deviation is stored in `importanceSD`. The coefficient of variation (i.e., mean/SD) is reported when `importance()` is called.

```
all.equal(importance(rf)[,'MeanDecreaseAccuracy'], ## True
  rf$importance[, 'MeanDecreaseAccuracy']/rf$importanceSD[, 'MeanDecreaseAccuracy'])
```

One can also get **local variable importance**, which is the mean decrease of accuracy of the predictors for each observation when it is OOB. To do this, set `localImp=T`. The results are stored in the component `localImportance`.

```
rf = randomForest(AHD ~ ., data=Heart2, localImp=T)
dim(rf$localImportance) ## 13 x 297
matplot(rf$localImportance[,1:5], type='l') ## first 5 observations
#rf$localImportance %>% rf$oob.times
```

Cross-validation to select m : The R caret package can do cross-validation to select the hyperparameter m . By default, a grid over the hyperparameter is selected by the `train()` function. To change that, use either `tuneLength=` to specify the number of values in a grid or `tuneGrid=` to specify the values in a data frame.

```
library(caret)
cvCtrl = trainControl(method="repeatedcv", number=5, repeats=4, ## 5-fold CV repeated 4 times
  classProbs=TRUE, summaryFunction=twoClassSummary)
fitRFCaret = train(x=Heart2[,1:13], y=Heart2$AHD, trControl=cvCtrl,
  #tuneGrid=data.frame(mtry=c(2,3,4,5)),
  #tuneLength=4,
  method="rf", metric="ROC", ntree=50) ## small ntree for class demonstration

fitRFCaret
names(fitRFCaret)
fitRFCaret$bestTune$mtry
fitRFCaret$finalModel$confusion ## same as table(Heart2$AHD, fitRFCaret$finalModel$predicted)
```

Again, the following generate “apparent error rates” that are too good.

```
table(Heart2$AHD, predict(fitRFCaret))
table(Heart2$AHD, predict(fitRFCaret, Heart2))
```

```
table(Heart2$AHD, predict(fitRFcaret$finalModel, Heart2))
```

`train(x, y)` is recommended because it does not create dummy variables for categorical predictors with more than two categories. If you use `train(formula, data)`, dummy variables are created for categorical predictors with more than two categories. This may cause problems when using functions such as `predict()`.

```
fitRFcaret2 = train(AHD ~ ., data=Heart2, trControl=cvCtrl,
                    method="rf", metric="ROC", ntree=20)
predict(fitRFcaret2$finalModel, Heart2) ## Error
fitRFcaret2$finalModel$importance
fitRFcaret$finalModel$importance
```

Example: Classification RF on the NCI60 dataset, which contains gene expression of 6830 genes for 64 tumor samples. The tumors are classified into 14 types.

```
library(ISLR)
str(NCI60)
table(NCI60$labs)
dd = data.frame(NCI60$data)
dd$labs = as.factor(NCI60$labs) ## randomForest() requires categorical outcome to be a factor
rf = randomForest(labs ~ ., data=dd, ntree=1000)
plot(rf$err.rate[,1], type='l', xlab='trees', ylab='Error')

## Remove those rare/unknown tumor types (those with only one obs)
idx = dd$labs %in% names(table(NCI60$labs))[table(NCI60$labs)>1]
dd2 = dd[idx,]
dd2$labs = as.factor(as.character(dd2$labs)) ## redefine the factor levels
rf2 = randomForest(labs ~ ., data=dd2, ntree=1000)
plot(rf2$err.rate[,1], type='l', xlab='trees', ylab='Error')
points(rf$err.rate[,1], type='l', col=2)
```

Example: Regression RF on the Wage dataset:

```
library(ISLR)
str(Wage)
levels(Wage$education)
Wage$education = factor(Wage$education, levels(Wage$education), ordered=T)

rf = randomForest(wage ~ .-logwage, data=Wage) ## all predictors except `logwage`
rf
names(rf)
rf$mtry; rf$ntree
```

The counterpart of `err.rate` is `mse`.

```
plot(rf) ## same as plot(rf$mse, type='l', xlab='trees', ylab='Error')
plot(Wage$wage, rf$predicted)
cor(Wage$wage, rf$predicted, method='spearman')

## Again, below is wrong.
plot(Wage$wage, predict(rf, Wage))
cor(Wage$wage, predict(rf, Wage), method='spearman')
```

For a regression model fitted with `randomForest()`, its component importance contains `IncNodePurity`. When `importance=T` is used, `%IncMSE` is the counterpart of `MeanDecreaseAccuracy` in classification RF.

```
rf = randomForest(wage ~ .-logwage, data=Wage, importance=T)
importance(rf)
varImpPlot(rf)
```

The smaller m the more spread-out of importance over the variables.

```
varImpPlot(randomForest(wage ~ .-logwage, data=Wage, mtry=1))
varImpPlot(randomForest(wage ~ .-logwage, data=Wage, mtry=3))
varImpPlot(randomForest(wage ~ .-logwage, data=Wage, mtry=5))
```

Individuals trees can be obtained with `getTree()`. For example, to get the 100th tree from the object `rf`, use `getTree(rf, 100)`. Unfortunately, the result of `getTree()` is in a format different from the frame format used by the `rpart` and `tree` packages, we cannot use the graphics functions from those packages. An R package `reprtree` posted on the GitHub can plot the individual trees. The plots are often not useful because the individual trees often have a high depth. To install the `reprtree` package, use `devtools::install_github('araastat/reprtree')`.

```
rf = randomForest(AHD ~ ., data=Heart2)
reprtree::plot.getTree(rf, k=100)
## redraw the tree without text
library(rpart)
tr100 = reprtree::as.tree(getTree(rf, k=100, labelVar=T), rf)
plot(tr100, type='uniform')
#text(tr100, split=T, cex=.8) ## adding text may make the plot look too crowded
```

The `reprtree` package also implements a method to generate “representative trees” for random forests (Banerjee, Ding, Noone (2012). Identifying representative trees from ensembles. *Statistics in Medicine*).

```
rep = reprtree::ReprTree(rf, Heart2) ## representative tree
reprtree::plot.reprtree(rep)
```

Notes on other R packages for bagging and random forests:

- `caret` ([document](#)) provides a unified interface (and a wrapper) for using `randomForest` and many other R packages. For example, it offers bagging of base models of type: `lda`, `pls`, `nb` (naive Bayes), `ctree`, `svm`, `nnet`. `names(getModelInfo())` provides a list of all trainable model types.
- `party` provides functions to build “conditional inference trees” (`ctree`) and to fit random forest models with conditional inference trees as base models.

8.1.3 Assignment

1. Reading for next lecture: ISLR 8.2, HOML 7.
2. ISLR Chapter 8 R Labs

8.2 Week 8 Day 2 (off for midterm)